

Krish Pradhan

Uni Id: 2359865

Workshop 2

```
Commands + Code + Text RAM Disk ^
[1] import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv("/content/drive/MyDrive/AI and Machine Learning/mnist_dataset.csv")

[3] from google.colab import drive
drive.mount('/content/drive')
Mounted at /content/drive

[6] print("Dataset preview:")
df.head()
```

Dataset preview:

	label	pixel_0	pixel_1	pixel_2	pixel_3	pixel_4	pixel_5	pixel_6	pixel_7	pixel_8	...	pixel_774	pixel_775	pixel_776	pixel_777	pixel_778	pixel_779	pixel_780	pixel_781	pi
0	5	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0
2	4	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0
3	1	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0
4	9	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0

5 rows × 785 columns

```
[7] print("Dataset info:")
print(df.info())

Dataset info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 60000 entries, 0 to 59999
Columns: 785 entries, label to pixel_783
dtypes: int64(785)
memory usage: 359.3 MB
None

[8] #extracting features (x) and target values (y)
x = df.iloc[:, 1:].values
y = df.iloc[:, 0].values

X.shape
(60000, 784)

[10] y.shape
(60000,)
```

```
01 plt.imshow(X[1].reshape(28, 28), cmap='gray')
    <matplotlib.image.AxesImage at 0x7971951d34d0>

[12] y[1]
    0

[13] X[0].shape
    (784,)
```

```
24 [14] from sklearn.model_selection import train_test_split
    from sklearn.preprocessing import OneHotEncoder

    # Splitting
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

05 [15] y_train.shape
    (48000,)

[16] # Reshape y_train and y_test to be column vectors
    y_train = y_train.reshape(-1, 1)
    y_test = y_test.reshape(-1, 1)

06 [17] y_train.shape
    (48000, 1)

    # Initialize the OneHotEncoder with correct parameter name
    encoder = OneHotEncoder(sparse_output=False, categories='auto')

    # Fit and transform the training labels, then transform the test labels
    y_train = encoder.fit_transform(y_train)
    y_test = encoder.transform(y_test)
```


```
07 [19] print("\nShapes:")
    print("X_train:", X_train.shape, "y_train:", y_train.shape)
    print("X_test:", X_test.shape, "y_test:", y_test.shape)

    Shapes:
    X_train: (48000, 784) y_train: (48000, 10)
    X_test: (12000, 784) y_test: (12000, 10)

08 [20] y_train[0]
    array([0., 0., 0., 0., 0., 1., 0., 0., 0., 0.])
```

▼ Some Helper Function:

▼ Softmax Function:

```
0s  import numpy as np

def softmax(z):
    """
    Compute the softmax probabilities for a given input matrix.

    Parameters:
    z (numpy.ndarray): Logits (raw scores) of shape (m, n), where
        - m is the number of samples.
        - n is the number of classes.

    Returns:
    numpy.ndarray: Softmax probability matrix of shape (m, n), where
        each row sums to 1 and represents the probability
        distribution over classes.

    Notes:
    - The input to softmax is typically computed as:  $z = XW + b$ .
    - Uses numerical stabilization by subtracting the max value per row.
    """

    # Your Code Here.
    z_shifted = z - np.max(z, axis=1, keepdims=True)
    exp_z = np.exp(z_shifted)
    softmax_probs = exp_z / np.sum(exp_z, axis=1, keepdims=True)
    return softmax_probs
```

▼ Softmax Test Case:


This test case checks that each row in the resulting softmax probabilities sums to 1, which is the fundamental property of softmax.

```
✓ [22] # Example test case
z_test = np.array([[2.0, 1.0, 0.1], [1.0, 1.0, 1.0]])
softmax_output = softmax(z_test)

# Verify if the sum of probabilities for each row is 1 using assert
row_sums = np.sum(softmax_output, axis=1)

# Assert that the sum of each row is 1
assert np.allclose(row_sums, 1), f"test failed: Row sums are {row_sums}"

print("Softmax function passed the test case!")
```

 Softmax function passed the test case!

```

def cross_entropy_loss(y_pred, y):
    """
    Compute the cross-entropy loss for a single sample.

    Parameters:
    y_pred (numpy.ndarray): Predicted probabilities of shape (c,) for a single sample,
                           where c is the number of classes.
    y (numpy.ndarray): True labels (one-hot encoded) of shape (c,), where c is the number of classes.

    Returns:
    float: Cross-entropy loss for the given sample.
    """

    # Avoid log(0) by adding a small value (epsilon)
    epsilon = 1e-15 # for not making (log(0))
    y_pred = np.clip(y_pred, epsilon, 1 - epsilon)

    # Calculate cross-entropy loss
    loss = -np.sum(y * np.log(y_pred))

    return loss

```

```

def cost_softmax(X, y, W, b):
    """
    Compute the softmax regression cost (cross-entropy loss).

    Parameters:
    X (numpy.ndarray): Feature matrix of shape (n, d), where n is the number of samples and d is the number of features.
    y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c), where c is the number of classes.
    W (numpy.ndarray): Weight matrix of shape (d, c).
    b (numpy.ndarray): Bias vector of shape (c,).

    Returns:
    float: The softmax cost (cross-entropy loss).
    """
    n = X.shape[0] # Number of samples
    z = np.dot(X, W) + b
    y_pred = softmax(z)
    cost = cross_entropy_loss(y_pred, y)
    return cost

```

```

[25] def compute_gradient_softmax(X, y, W, b):
    """
    Compute the gradients of the cost function with respect to weights and biases.

    Parameters:
    X (numpy.ndarray): Feature matrix of shape (n, d).
    y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c).
    W (numpy.ndarray): Weight matrix of shape (d, c).
    b (numpy.ndarray): Bias vector of shape (c,).

    Returns:
    tuple: Gradients with respect to weights (d, c) and biases (c,).
    """
    n, d = X.shape #(total_number_of_row, total_columns)
    z = np.dot(X, W) + b
    y_pred = softmax(z)

```

```

def compute_gradient_softmax(X, y, W, b):
    """
    Compute the gradients of the cost function with respect to weights and biases.

    Parameters:
    X (numpy.ndarray): Feature matrix of shape (n, d).
    y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c).
    W (numpy.ndarray): Weight matrix of shape (d, c).
    b (numpy.ndarray): Bias vector of shape (c,).

    Returns:
    tuple: Gradients with respect to weights (d, c) and biases (c,).
    """
    n, d = X.shape #(total_number_of_row, total_columns)
    z = np.dot(X, W) + b
    y_pred = softmax(z)

    grad_W = np.dot(X.T, (y_pred - y)) / n # Gradient with respect to weights
    grad_b = np.sum(y_pred - y, axis=0) / n # Gradient with respect to biases

    return grad_W, grad_b

```

```
[26] def gradient_descent_softmax(X, y, W, b, alpha, n_iter, show_cost=False):
    """
    Perform gradient descent to optimize the weights and biases.

    Parameters:
    X (numpy.ndarray): Feature matrix of shape (n, d).
    y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c).
    W (numpy.ndarray): Weight matrix of shape (d, c).
    b (numpy.ndarray): Bias vector of shape (c,).
    alpha (float): Learning rate.
    n_iter (int): Number of iterations.
    show_cost (bool): Whether to display the cost at intervals.

    Returns:
    tuple: Optimized weights, biases, and cost history.
    """
    cost_history = []

    for i in range(n_iter):
        # Compute gradients
        grad_W, grad_b = compute_gradient_softmax(X, y, W, b)

        # Update weights and biases using the gradients
        W -= alpha * grad_W
        b -= alpha * grad_b

        # Compute and store cost
        cost = cost_softmax(X, y, W, b)
        cost_history.append(cost)

        # Print cost at regular intervals
        if show_cost and (i % 100 == 0 or i == n_iter - 1):
            print(f"Iteration {i}: Cost = {cost:.6f}")

    return W, b, cost_history
```

▼ Prediction Function:


```
[27] def predict_softmax(X, W, b):  
    """  
    Predict the class labels for a set of samples using the trained softmax model.  
  
    Parameters:  
    X (numpy.ndarray): Feature matrix of shape (n, d), where n is the number of samples and d is the number of features.  
    W (numpy.ndarray): Weight matrix of shape (d, c), where c is the number of classes.  
    b (numpy.ndarray): Bias vector of shape (c,).  
  
    Returns:  
    numpy.ndarray: Predicted class labels of shape (n,), where each value is the index of the predicted class.  
    """  
    z = np.dot(X, W) + b # Compute the scores (logits)  
    y_pred = softmax(z) # Get the probabilities using the softmax function  
  
    # Assign the class with the highest probability  
    predicted_classes = np.argmax(y_pred, axis=1)  
  
    return predicted_classes
```

```
[28] y_train.shape
```

(48000, 10)

```
[29] # Initialize the weights and biases  
d = X_train.shape[1] # Number of features  
c = y_train.shape[1] # Number of classes  
W = np.random.randn(d, c) * 0.01 # Small random weights  
b = np.zeros(c) # Bias initialized to 0  
  
# Set hyperparameters  
alpha = 0.1 # Learning rate  
n_iter = 1000 # Number of iterations  
  
# Train the model using gradient descent  
W_opt, b_opt, cost_history = gradient_descent_softmax(X_train, y_train, W, b, alpha, n_iter, show_cost=True)  
  
# Plot the cost history to visualize the convergence  
plt.plot(cost_history)  
plt.title('Cost Function vs. Iterations')  
plt.xlabel('Iterations')  
plt.ylabel('Cost')  
plt.grid(True)  
plt.show()  
  
# Predict on the test set  
y_pred_test = predict_softmax(X_test, W_opt, b_opt)  
  
# Evaluate accuracy  
y_test_labels = np.argmax(y_test, axis=1) # True labels in numeric form  
#accuracy = np.mean(y_pred_test == y_test_labels)  
#print(f"Test accuracy: {accuracy * 100:.2f}%")
```

```
Iteration 0: Cost = 1149736.612632  
Iteration 100: Cost = 582747.027293  
Iteration 200: Cost = 143895.863400  
Iteration 300: Cost = 137412.689552  
Iteration 400: Cost = 218113.846292  
Iteration 500: Cost = 137911.223933  
Iteration 600: Cost = 133376.355944  
Iteration 700: Cost = 157737.261412  
Iteration 800: Cost = 138856.799485  
Iteration 900: Cost = 167943.940920  
Iteration 999: Cost = 240358.253716
```




```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score

# Evaluation Function
def evaluate_classification(y_true, y_pred):
    """
    Evaluate classification performance using confusion matrix, precision, recall, and F1-score.

    Parameters:
    y_true (numpy.ndarray): True labels
    y_pred (numpy.ndarray): Predicted labels

    Returns:
    tuple: Confusion matrix, precision, recall, F1 score
    """
    # Compute confusion matrix
    cm = confusion_matrix(y_true, y_pred)

    # Compute precision, recall, and F1-score
    precision = precision_score(y_true, y_pred, average='weighted')
    recall = recall_score(y_true, y_pred, average='weighted')
    f1 = f1_score(y_true, y_pred, average='weighted')

    return cm, precision, recall, f1

# Predict on the test set
y_pred_test = predict_softmax(X_test, W_opt, b_opt)
# Evaluate accuracy
y_test_labels = np.argmax(y_test, axis=1) # True labels in numeric form

# Evaluate the model
cm, precision, recall, f1 = evaluate_classification(y_test_labels, y_pred_test)

# Print the evaluation metrics
print("\nConfusion Matrix:")
print(cm)
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"F1-Score: {f1:.2f}")

# Visualizing the Confusion Matrix
fig, ax = plt.subplots(figsize=(6, 6))
cax = ax.imshow(cm, cmap='Blues') # Use a color map for better visualization

# Set tick labels for the axes
ax.set_xticks(range(3))
ax.set_yticks(range(3))
ax.set_xticklabels([f'Predicted {i}' for i in range(3)])
ax.set_yticklabels([f'Actual {i}' for i in range(3)])

# Add labels to each cell in the confusion matrix
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        ax.text(j, i, cm[i, j], ha='center', va='center', color='white' if cm[i, j] > np.max(cm) / 2 else 'black')

# Add grid lines and axis labels
ax.grid(False)
plt.title('Confusion Matrix', fontsize=14)

```

```

# Add labels to each cell in the confusion matrix
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        ax.text(j, i, cm[i, j], ha='center', va='center', color='white' if cm[i, j] > np.max(cm) / 2 else 'black')

# Add grid lines and axis labels
ax.grid(False)
plt.title('Confusion Matrix', fontsize=14)
plt.xlabel('Predicted Label', fontsize=12)
plt.ylabel('Actual Label', fontsize=12)

# Adjust layout
plt.tight_layout()
plt.colorbar(cax)
plt.show()

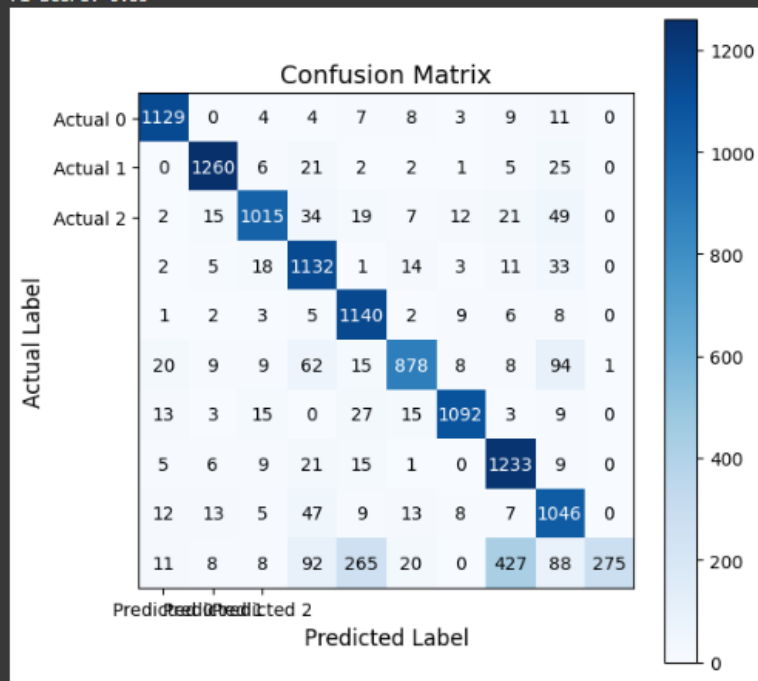
```



```


Confusion Matrix:
[[1129  0  4  7  8  3  9 11  0]
 [  0 1260  6 21  2  2  1  5 25  0]
 [  2  15 1015 34 19  7 12 21 49  0]
 [  2  5  18 1132  1 14  3 11 33  0]
 [  1  2  3  5 1140  2  9  6  8  0]
 [ 20  9  9  62 15 878  8  8 94  1]
 [ 13  3 15  0 27 15 1092  3  9  0]
 [  5  6  9 21 15  1  0 1233  9  0]
 [ 12 13  5 47  9 13  8  7 1046  0]
 [ 11  8  8 92 265 20  0 427 88 275]]
Precision: 0.87
Recall: 0.85
F1-Score: 0.83

```



```
✓ [33] # Reshape the first image to (28, 28)
0s image = X_train[0].reshape(28, 28)

# Display the image using plt.imshow
plt.imshow(image, cmap='gray') # 'gray' colormap for grayscale images
plt.show()
```



```
✓ y_train[0]
```

```
0s array([0., 0., 0., 0., 0., 1., 0., 0., 0., 0.])
```

```
✓ [35]
0s single_image = X_train[0].reshape(1, -1) # Reshape to (1, 784) as it needs to be 2D

predicted_class = predict_softmax(single_image, W_opt, b_opt)
```

```
✓ [36] print("Predicted class for the single image:", predicted_class[0])
0s
➡ Predicted class for the single image: 5
```